

One frozen model, every capability: fastrho vs the incumbents

fastrho	yes	0.88	0.77	yes	yes	yes	1×
pyrho	yes	0.80	0.68	no	no (table)	no (SMC++)	70×
ReLERNN	no (windows)	0.41	0.26	bootstrap	no (retrain)	trained-in	8527×
	Per-interval resolution	Fine-scale r (const $n=20$, 25kb)	Hotspot AUPRC (const $n=20$)	1-pass calibrated uncertainty	No per-dataset table / retrain	No demography required	Relative speed